

Best Practices to Build Effective Agentic AI Systems

Created & Narrated by:

Dipanjan Sarkar

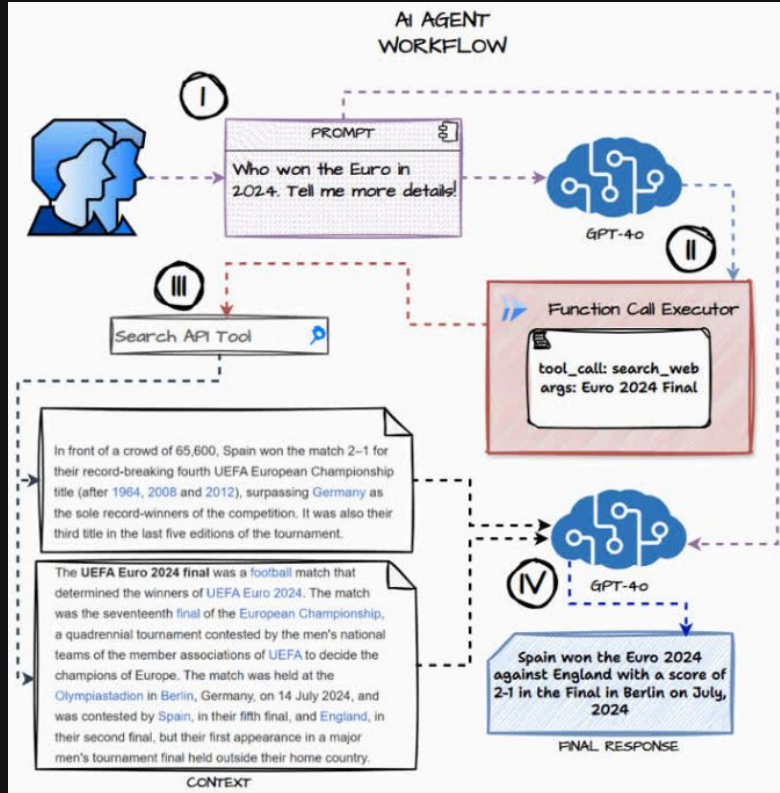
Head of Community & Principal AI Scientist @ Analytics Vidhya

Google Developer Expert - ML & Cloud Champion Innovator

Published Author

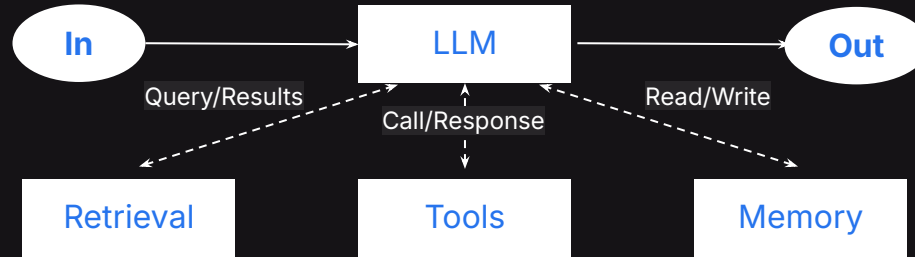


Recap: What is an Agentic AI System



- An **Agentic AI System** is usually an autonomous system that operates independently over extended periods, using various tools and flows to accomplish complex tasks.
- Agentic AI Systems can be further categorized as:
 - **Workflows** are systems where LLMs and tools are orchestrated through predefined paths.
 - **Agents**, on the other hand, are systems where LLMs dynamically direct their own processes and tool usage, maintaining control over how they accomplish tasks.

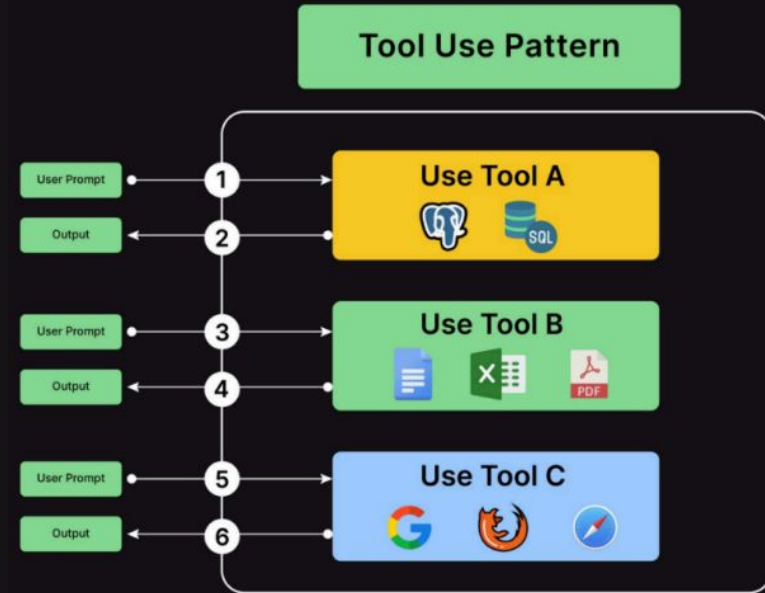
1. Build the Key Components for your Agent



- The basic building block of Agentic AI Systems is an **LLM enhanced with augmentations such as retrieval, tools, and memory**.
- Powerful LLM platforms have these in-built. When using APIs you would need to **connect the LLM with relevant tools, memory and databases** so that they can generate their own search queries, select appropriate tools, and determine what information to retain.
- Anthropic recommends focusing on **two** key aspects of the implementation:
 - Tailoring these capabilities to your specific use case.
 - Ensuring they provide an easy, well-documented interface for your LLM.

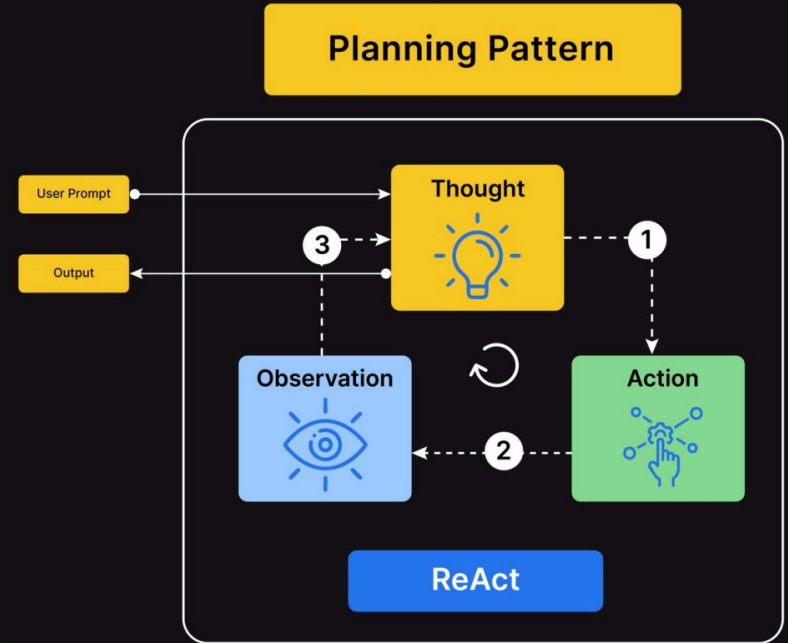
2. Start with Tool-Use Single ReAct Agents

- Most **ReAct Tool-Use Agents** already have **planning capabilities**.
- These systems can easily handle **10-15 tools** easily.
- They can also handle multi-step and multi-tool call executions easily.
- Key drivers here include:
 - Well-defined tool schemas for accurate function calling.
 - Well-structured system prompt with detailed instructions.
 - Powerful LLMs already trained for function (tool) calling.
- Do NOT jump into multi-agent systems immediately - they are notoriously hard to control and debug regardless of present advancements.

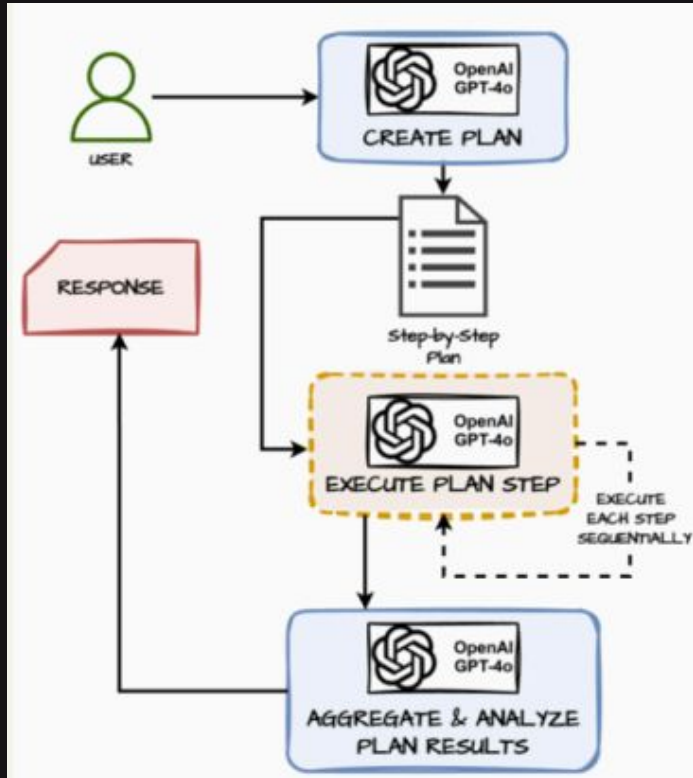


3. Planning Agents for Complex Task Execution

- Most **ReAct Agents** already have **planning built-in** so first start with simple ReAct Agents.
- If tasks are more complex and require explicit planning consider adding in additional planning modules in the Agent.
- Planning modules or patterns are typically:
 - **Static Planners** with Parallel Task Execution & Synthesis
 - **Dynamic Planners** with Task Execution, Reflection & Replanning
- Do **NOT** add extra planning steps or modules unless absolutely necessary as this **increases system latency**
- It is especially useful in complex reasoning, inference and validation scenarios.



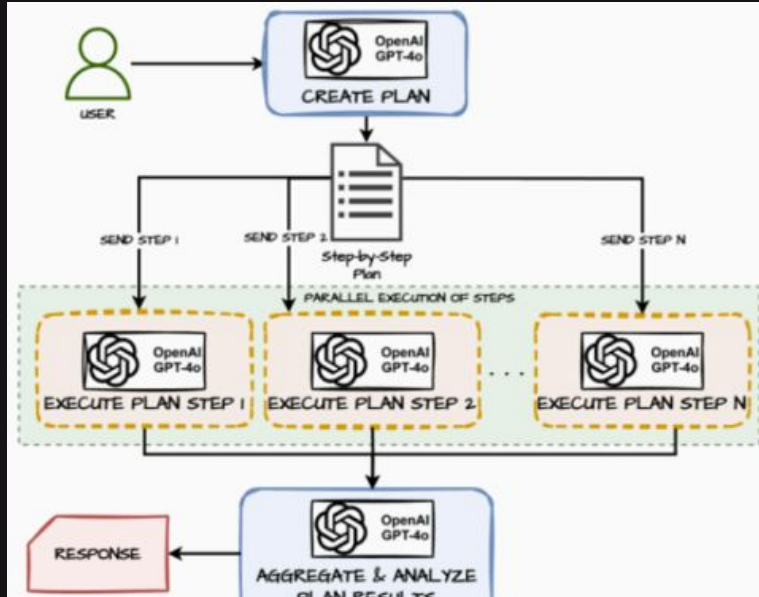
3. Planning Agents for Complex Task Execution



Dynamic Planners

- Use planning to break down a task into multiple steps.
- Executes one step at a time.
- Reflects on results of steps already executed.
- Uses reflection to replan remaining steps (if needed)
- Repeats till all steps are executed.
- Synthesizes results from all steps and generates final response.
- Useful when tasks may have dependencies among each other.

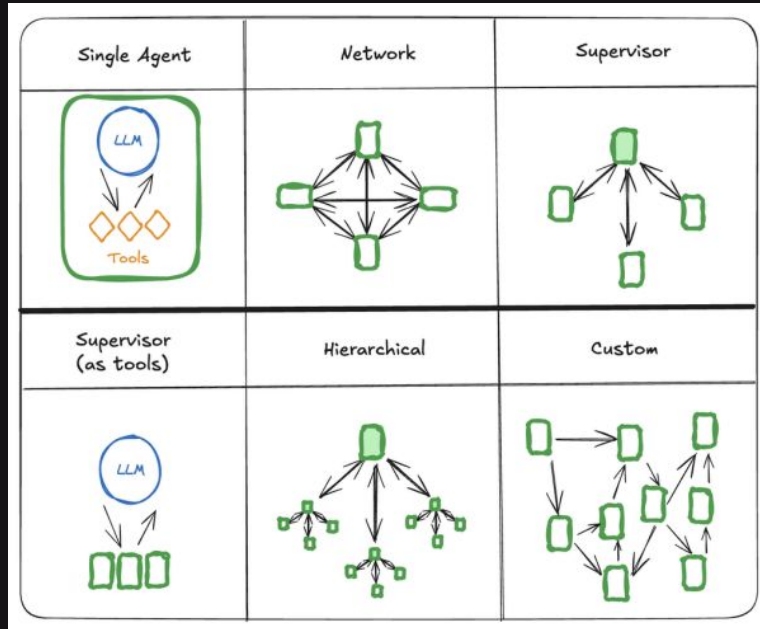
3. Planning Agents for Complex Task Execution



Static Planners


- Use planning to break down a task into multiple steps.
- Execute all steps in parallel
- Synthesize results from all steps and generate final response (map-reduce)
- Useful when steps do not have dependencies

4. Multi-Agent Systems for Efficient Task Distribution




- **Network:** Each agent can communicate with every other agent. Any agent can decide which other agent to call next
- **Supervisor:** Each agent communicates with a single supervisor agent. Supervisor agent makes decisions on which agent should be called next
- **Hierarchical:** Multi-agent system with a supervisor of supervisors. This is a generalization of the supervisor architecture and allows for more complex control flows
- Always start with simple supervisor or network architecture and then expand.
- Create separate agents based on specific processes, tasks and flows.

5. Reflection for Critiquing & Improvements



- Agents using reflection will leverage **one LLM call to generate a response** while **another provides evaluation and feedback in a loop**.
- This workflow is particularly effective when we have **clear evaluation criteria, and when iterative refinement provides measurable value**.
- **The two signs of good fit are:**
 - LLM responses can be demonstrably improved when a human articulates their feedback
 - The LLM can provide such feedback.

5. Reflection for Critiquing & Improvements



- This is analogous to the iterative writing process a human writer might go through when producing a polished document or when developed code is reviewed, tested and then improved.
- **Examples where reflection is useful:**
 - **Improving quality of RAG** retrieval and responses
 - **Judging, grading and critiquing** the quality of an LLM response
 - **Validating** specific criteria and guidelines e.g claims processing

Summary of Key Takeaways

- 1 Start by building key components of your Agentic AI System.
- 2 Start with Simple Tool-Use ReAct Single Agent Systems.
- 3 Add Planning only if the above fails and you need to add explicit planning modules for complex tasks.
- 4 Add Reflection for handling any tasks around grading, critiquing, improving responses at any step in the agent.
- 5 Add Routing to handle multiple flows or agents reliably.
- 6 Consider Multi-Agent systems when you have multiple processes, workflows, too many tools to handle and you can segregate tasks to specific agents with a set of tools.
- 7 Do not add in too many Agents unless absolutely necessary.
- 8 Monitoring and Debugging is super useful to check for common failure patterns.

Thanks!